

Quartz.Net Scheduler in Depth

DRAFT

Introduction

What is a Job Scheduler?

Wikipedia defines a job scheduler as:

“A job scheduler is a software application that is in charge of unattended background executions, commonly known for historical reasons as batch processing.”

If we paraphrase this definition a little, we can say that a job scheduler is an application that lets us do unattended job processing.

Job schedulers are all around us, since they are included in many systems that we use every day. Here are a few examples of systems that provide some sort of job scheduler:

- Operating systems (task scheduler in windows or cron in linux)
- Databases (SQL Agent in SQL Server)

What is Quartz.Net?

Quartz .Net is an open source job scheduler for the .Net platform. It has been ported to run on the .Net framework from the original Quartz Java version. Quartz.Net is written in C# and is a straight code port, with the Java code being basically translated to C#. The Quartz.Net scheduler can be run as a standalone scheduler or it can also be embedded in your application.

Development of Quartz.Net 1.0 was originally done on Sourceforge, and the latest build for the 1.0 version can still be downloaded from there (<http://sourceforge.net/projects/quartznet/files/quartznet/>). However, for version 2.0, development of Quartz.Net has moved to Github (<https://github.com/quartznet/quartznet>). The binaries are still available on Sourceforge. You can download the latest binary release of Quartz.Net 2.0 from <http://sourceforge.net/projects/quartznet/files/quartznet/>

The core Quartz.Net library is also available as a NuGet package and you can visit the Quartz.Net page on NuGet here: <http://nuget.org/packages/Quartz>.

Why Use Quartz.Net?

Operating systems and databases all have some sort of task scheduling capability. Why should you consider using Quartz instead of these built-in schedulers? Here are some reasons:

- You want to limit the total number of jobs that run at the same time
- You need more than one server to run all the jobs
- You want a central point of control for all jobs and servers

System Requirements

- .Net Framework 3.5 or higher
- Windows XP SP3 or higher operating system

- Visual Studio 2008 or 2010

Part 1 – Quick Start

This quick start guide is designed to get you up and running as fast as possible. To do this we will skip most of the details and focus on:

- downloading and installing Quartz.Net
- downloading the sample code
- getting a Quartz.Net server instance configured and running
- scheduling and running a built-in job
- creating a simple custom job
- scheduling and running the custom job

Downloading and Setting Up Quartz.Net

Download the code samples from <https://github.com/jvilalta/QuartzNetBook/downloads>. You can unzip the file to any folder on your hard drive, but we will be using `c:\QuartzNetBook` throughout the book. Once unzipped, your folder structure underneath the installation directory `c:\QuartzNetBook` will look like this:




Figure 1

Find the server folder and open a command prompt in that folder.




Figure 2

To run the scheduler, simply execute Quartz.Server.exe by typing Quartz.Server.exe in the command window.




Figure 3

Once you execute the command, the scheduler will start running. If your operating system brings up a prompt about the windows firewall, click on unblock so that you can connect to the scheduler using the built-in remoting support (we will take a look at this later on).

As the scheduler is starting, you will see several informational messages scroll by in the command prompt window. The last message on the screen should say “Running, press Control+C to exit”.

```
C:\Windows\system32\cmd.exe - Quartz.Server.exe
2011-12-08 18:14:58.580 [6136] INFO    Topshelf.Model.ServiceCoordinator.OnService
Event(<0>) - (quartz.server) Created
2011-12-08 18:14:58.621 [6372] INFO    Topshelf.Model.ServiceCoordinator.OnService
Event(<0>) - (quartz.server) Starting
2011-12-08 18:14:58.647 [6136] INFO    Quartz.Xml.XMLSchedulingDataProcessor.Proce
ssFile(C:\git\jvilalta\quartznet\src\Quartz\Xml\xml\SchedulingDataProcessor.cs:18
1) - Parsing XML file: C:\QuartzNetBook\setup\quartz_jobs.xml with systemId: ~/q
uartz_jobs.xml
2011-12-08 18:14:59.015 [6136] INFO    Quartz.Xml.XMLSchedulingDataProcessor.Sched
uleJobs(C:\git\jvilalta\quartznet\src\Quartz\Xml\xml\SchedulingDataProcessor.cs:6
72) - Adding 0 jobs, 0 triggers.
2011-12-08 18:14:59.021 [6136] INFO    Quartz.Core.QuartzScheduler.Start(C:\git\jv
ilalta\quartznet\src\Quartz\Core\QuartzScheduler.cs:443) - Scheduler ServerSched
uler $ NON_CLUSTERED started.
2011-12-08 18:15:02.027 [6136] INFO    Topshelf.Builders.LocalServiceBuilder`1.Sta
rtService(<0>) - Scheduler started successfully
2011-12-08 18:15:02.027 [6136] INFO    Topshelf.Model.LocalServiceController`1.Cal
lAction(<0>) - [quartz.server] Start complete
2011-12-08 18:15:02.029 [8032] INFO    Topshelf.Model.ServiceCoordinator.OnService
Event(<0>) - (quartz.server) Running
2011-12-08 18:15:03.031 [728] INFO     Topshelf.Hosts.ConsoleRunHost.Run(<0>) - [Top
shelf] After starting services complete
2011-12-08 18:15:03.032 [728] INFO     Topshelf.Hosts.ConsoleRunHost.Run(<0>) - [Top
shelf] Running, press Control+C to exit.
```

Figure 4

To stop the scheduler, press Control-C, and watch as the scheduler shuts down. Once the scheduler completes the shutdown process, you should be back at the command prompt, and the last line outputted by the scheduler will say “Stopped”.

```
C:\Windows\system32\cmd.exe
eduler $ NON_CLUSTERED paused.
2011-12-19 12:58:06.400 [4416] INFO    Quartz.Simpl.RemotingSchedulerExporter.UnBi
nd(C:\git\jvilalta\quartznet\src\Quartz\Simpl\RemotingSchedulerExporter.cs:158)
- Successfully disconnected remutable scheduler
2011-12-19 12:58:06.404 [4416] INFO    Quartz.Core.QuartzScheduler.Shutdown(C:\git
\jvilalta\quartznet\src\Quartz\Core\QuartzScheduler.cs:652) - Scheduler ServerSc
heduler $ NON_CLUSTERED Shutdown complete.
2011-12-19 12:58:06.404 [4416] INFO    Topshelf.Builders.LocalServiceBuilder`1.Sto
pService(<0>) - Scheduler shutdown complete
2011-12-19 12:58:06.405 [4416] INFO    Topshelf.Model.LocalServiceController`1.Cal
lAction(<0>) - [quartz.server] Stop complete
2011-12-19 12:58:06.406 [6564] INFO    Topshelf.Model.ServiceCoordinator.OnService
Event(<0>) - (quartz.server) Stopped
2011-12-19 12:58:06.452 [4416] INFO    Topshelf.Model.ServiceCoordinator.OnService
Event(<0>) - (quartz.server) Unloading
2011-12-19 12:58:06.461 [6564] INFO    Topshelf.Model.LocalServiceController`1.Cal
lAction(<0>) - [quartz.server] Unload complete
2011-12-19 12:58:06.462 [580] INFO     Topshelf.Model.ServiceCoordinator.OnServiceE
vent(<0>) - (quartz.server) Unloaded
2011-12-19 12:58:07.463 [8184] INFO    Topshelf.Hosts.ConsoleRunHost.ShutdownCoord
inator(<0>) - [Topshelf] After stopping services complete
2011-12-19 12:58:07.498 [8184] INFO    Topshelf.Hosts.ConsoleRunHost.ShutdownCoord
inator(<0>) - [Topshelf] Stopped
C:\QuartzNetBook\server>
```

Figure 5

Leave the scheduler shut down for now, as we will now take a look at how to configure the scheduler. Keep in mind as you work through the configuration examples that any configuration changes made while the scheduler is running will not be applied to the scheduler until after it is restarted.

A Quick Look at Configuration

For the quick start we'll look at the simplest way of configuring Quartz.Net. The default setup of Quartz.Net includes a quartz.config file that holds the configuration parameters for the scheduler. Let's take a look at the sample configuration file.

```
# You can configure your scheduler in either <quartz> configuration section
# or in quartz properties file
# Configuration section has precedence

quartz.scheduler.instanceName = ServerScheduler

# configure thread pool info
quartz.threadPool.type = Quartz.Simpl.SimpleThreadPool, Quartz
quartz.threadPool.threadCount = 10
quartz.threadPool.threadPriority = Normal

# job initialization plugin handles our xml reading, without it defaults are
used
quartz.plugin.xml.type = Quartz.Plugin.Xml.XMLSchedulingDataProcessorPlugin,
Quartz
quartz.plugin.xml.fileNames = ~/quartz_jobs.xml

# export this server to remoting context
quartz.scheduler.exporter.type = Quartz.Simpl.RemotingSchedulerExporter,
Quartz
quartz.scheduler.exporter.port = 555
quartz.scheduler.exporter.bindName = QuartzScheduler
quartz.scheduler.exporter.channelType = tcp
quartz.scheduler.exporter.channelName = httpQuartz
```

This is the standard configuration that comes with the Quartz.Net distribution. It creates a scheduler that can run up to 10 jobs at a time. It also tells the scheduler factory that it should load the jobs described in the quartz_jobs.xml file into the scheduler. If we wanted to be able to run more than 10 jobs at a time, then we can change the value of quartz.threadPool.threadCount to say 15. This would tell the scheduler to execute up to 15 jobs at a time. If you wanted to load jobs into the scheduler from a file other than the default, all you have to do is change the value of quartz.plugin.xml.fileNames. We'll discuss configuration in depth later in the book, and in the following section you'll get a chance to schedule a job using the quartz_jobs.xml file.

Running an Executable or Batch File Using the Built-in Native Job

Quartz.Net provides a built-in job that lets you execute any file that can be run from a command prompt. The built-in job that allows you to do this is called the NativeJob. Let's take a look at how we can execute a batch file using the NativeJob. To illustrate the use of this job type, we will be running a batch file that writes the current date and time to a text file.

Configuring the Scheduler

First, we will configure the Quartz.Net scheduler so that it knows that it must execute the job that runs our batch file. To do this, we will be using the default job configuration mechanism, which is simply an xml file where we include all the jobs that the scheduler needs to run.

Here is what a barebones quartz_jobs.xml file looks like:

```
<?xml version="1.0" encoding="UTF-8"?>
<!-- This file contains job definitions in schema version 2.0 format -->
<job-scheduling-data
  xmlns="http://quartznet.sourceforge.net/JobSchedulingData"
  xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance" version="2.0">
  <processing-directives>
    <overwrite-existing-data>true</overwrite-existing-data>
  </processing-directives>
  <schedule>
  </schedule>
</job-scheduling-data>
```

We will add our job configuration data inside the `<schedule>` element in the above file. This is what the above file will look like after we add the configuration information to tell the scheduler to run our batch file:

```
<?xml version="1.0" encoding="UTF-8"?>
<!-- This file contains job definitions in schema version 2.0 format -->
<job-scheduling-data
  xmlns="http://quartznet.sourceforge.net/JobSchedulingData"
  xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance" version="2.0">
  <processing-directives>
    <overwrite-existing-data>true</overwrite-existing-data>
  </processing-directives>
  <schedule>
    <job>
      <name>nativeJobExample</name>
      <group>nativeJobExampleGroup</group>
      <description>Sample job for Quartz Server</description>
      <job-type>Quartz.Job.NativeJob, Quartz</job-type>
      <job-data-map>
        <entry>
          <key>command</key>
          <value>native_job_example.bat</value>
        </entry>
        <entry>
          <key>consumeStreams</key>
          <value>true</value>
        </entry>
      </job-data-map>
    </job>
    <trigger>
      <simple>
        <name>nativeJobExampleSimpleTrigger</name>
        <group>nativeJobExampleSimpleTriggerGroup</group>
        <description>Simple trigger to simply fire sample job</description>
```



```

        <job-name>nativeJobExample</job-name>
        <job-group>nativeJobExampleGroup</job-group>
        <misfire-instruction>SmartPolicy</misfire-instruction>
        <repeat-count>5</repeat-count>
        <repeat-interval>10000</repeat-interval>
    </simple>
</trigger>
</schedule>
</job-scheduling-data>

```

If you'd like to follow along this explanation, all of the files we are examining are located in the `c:\QuartzNetBook\samples\quickstart` folder.

As you can see, we've added two new elements under `<schedule>`: a `<job>` element and a `<trigger>` element. In order to get a job to execute under Quartz.Net, a trigger must be attached to it. The job describes the work to be done and the trigger tells the scheduler when this job should execute.

The <job> Element

Let us take a closer look at the `<job>` section now. This section describes the job that we want to execute in detail.

First, we give the job a name using the `<name>` element. Then we assign it to a group using the `<group>` element. We will highlight later on why these two items are important, but for now think of them as uniquely identifying a job.

Next in the `<job>` section is the `<description>` element, which describes in a friendly format what the job does. The `<job-type>` element tells the scheduler the name of the .NET class that must be instantiated to execute this job.

After the `<job-type>` element comes the `<job-data-map>` element.

The <job-data-map> Element

The `<job-data-map>` element is a container for all of the name-value pairs of data that you want to pass in to the job. These name value pairs become part of the job context during execution, thus allowing you to use the same job type but have it execute with different parameters. In this example we are passing in the name of the batch file that we want to execute. If we wanted to configure two different jobs different batch files to run, then we only need to change the name of the file that we are passing in to the job.

Inside the `<job-data-map>` we find one or more `<entry>` elements, which correspond to each name-value pair. Each `<entry>` element will have a `<key>` element and a `<value>` element. As you can tell from the names, the key element is the string that will be used to look up the value once inside the job. The value element is the actual value that is returned from the job data map.

In our example, the key is “command”, which tells that job that the value is the name of the batch file that we want to run. The native job supports other parameters, but we’ll skip those since they’re not important to running our example.

The <trigger> Element

It’s time now to take a deeper look at the <trigger> element. We will look at how to define and add triggers to our job. Triggers are responsible for kicking off jobs.

First, let’s look at how we define triggers. The Quartz.Net distribution includes two types of triggers: simple triggers and cron triggers. In our example we’ll use a simple trigger. A simple trigger is a trigger that fires a certain number of times with a certain delay between each firing. We’ll discuss cron triggers later.

In order to define a simple trigger we must start off with the <simple> tag, which serves as a container for the simple trigger. Next, we give our trigger a name using the <name> element and assign it to a group using the <group> element. You can also specify a description for the trigger using the <description> element. This is all the same as what we did for the job itself. These elements are also used for common across all trigger types, including cron triggers.

The <repeat-count> element will define how many times a simple trigger will repeat. The <repeat-interval> element specifies how much time (in milliseconds) will pass between trigger firings. These two elements together are what in essence define a simple trigger.

Sometimes the scheduler is busy running jobs, so some triggers might not be able to fire a job. When this happens, it’s called a misfire. The <misfire-instruction> element tells the scheduler what to do. For our example the misfire policy isn’t important, so we’ll skip it for now and will talk about them in depth later.

Linking the Trigger to the Job

Finally, we must link the job and the trigger together so that we can get our job to fire. This is done by setting the <job-name> element to the same value that we set the job’s <name> element and by setting the <job-group> to the same value that we set the job’s <group> element. These two values uniquely identify a job, so both of them must be provided in order to assign a trigger to a job.

Executing the Batch File

Now let’s take a quick look at the batch file before we execute it using the scheduler. First, copy the native_job_example.bat and quartz_jobs.xml files from c:\QuartzNetBook\samples\quickstart to c:\QuartzNetBook\server. There is a quartz_jobs.xml file already there, so when prompted to ask if you want to overwrite the file, go ahead and select the option to replace the file.

Before we run the batch file, take a look at what it does. Right click on the native_job_example.bat file and select edit. You will see that all it does is write the current date and time to a text file called run_times.txt. It’s a pretty simple file but it helps us to visualize what is going on behind the scenes when the scheduler is running. Double click on the file to execute it. If the file runs successfully you will

find a file named `run_times.txt` in the same directory (`c:\QuartzNetBook\server`). Open the file by double clicking on it. The file should contain one line with the current date and time.

It's time now to have the scheduler run the file. We'll do that from the command prompt so that we can see what the scheduler is doing. Open a command prompt at `c:\QuartzNetBook\server` and start the scheduler by typing `Quartz.Server.exe` and pressing Enter.




Figure 6

If all goes well the scheduler should start up again like before. However, now it will run our batch file every 10 seconds. You will know that the job is running because the scheduler will keep writing output to the command prompt. The job is configured to run 5 times, so after about a minute, the job will complete its execution and the scheduler will stop writing to the command prompt.

Your command prompt will look something like this:

Figure 7

Press Control-C to stop the scheduler. The scheduler should stop and take you back to the prompt.

Figure 8

What should have happened as the scheduler ran is that the batch file should have been called 6 times, every 10 seconds. To verify this, open the run_times.txt file again. You should see 6 date time entries in the file, about 10 seconds apart. But wait, we set the repeat count to 5, not 6, so why are there 6 entries instead of 5? Well, this is because the repeat count does not include the first run. If you wanted the job

to run 5 times, you would set the repeat count to 4 and this way you would get one initial run plus 4 repeat runs.

Creating a Custom Job

Now that we have scheduled one of the built-in jobs, it's time to create our own custom job. In order to create a custom job, we need to implement the IJob interface. This is a very simple interface, since it only requires us to implement one method. Here's the IJob interface definition:

```
public interface IJob
{
    void Execute(IJobExecutionContext context);
}
```

As you can see, we only need to implement the Execute method. In this method is where our job's code to "do something" will be written.

The QuickStartJob

The custom job we will implement will be a simple one. This job will do the same work as our previous batch file: it will write the current date and time to a text file. First, let's open the QuartzNetBookSamples.sln solution file using Visual Studio 2010. Once in Visual Studio, open the QuickStart project. In it you will find the QuickStartJob.cs file. This is the file that holds the code for our simple, yet custom job.

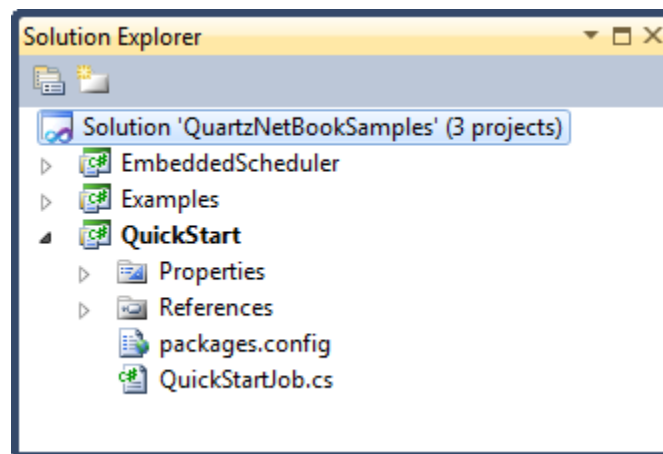


Figure 9

Here is the code for the QuickStartJob.cs file:

```
using System;
using System.IO;
using Quartz;

namespace QuickStart
{
    public class QuickStartJob : IJob
    {

```

```

public void Execute(IJobExecutionContext context)
{
    var fileName = "QuickStartJob_runs.txt";
    using (StreamWriter writer =
        new StreamWriter(File.Open(fileName, FileMode.Append)))
    {
        writer.WriteLine(DateTime.Now.ToString());
        writer.Flush();
        writer.Close();
    }
}
}
}

```

As you can see, we have implemented the Execute method with what would be the equivalent functionality of our batch file.

Now let's update the quartz_jobs.xml file so that in addition to running the batch file, the scheduler also runs our custom job.

Updating the Job Configuration File (quartz_jobs.xml)

In order to get the scheduler to run our custom job, we need to add it to the quartz_jobs.xml. We also need to add a trigger so that the job gets executed, but this is a fairly straightforward process.

First, we'll add the job by adding the following xml snippet just before the </schedule> tag:

```

<job>
  <name>quickStartJobExample</name>
  <group>quickStartJobExampleGroup</group>
  <description>My first Quartz.Net job</description>
  <job-type>QuickStart.QuickStartJob, QuickStart</job-type>
</job>

```

Before we move on to adding the trigger, I'd like to highlight the fact that since we have created a new job in a separate dll, the <job-type> element must reference that new assembly.

Now, let's add the trigger. We're going to be adding the same kind of trigger as in the previous example, but we'll have to reference the new job we've created. Underneath the job element that we've just added, let's add the following trigger element:

```

<trigger>
  <simple>
    <name>quickStartJobExampleSimpleTrigger</name>
    <group>quickStartJobExampleSimpleTriggerGroup</group>
    <description>Trigger that fires the QuickStartJob</description>
    <job-name>quickStartJobExample</job-name>
    <job-group>quickStartJobExampleGroup</job-group>
    <misfire-instruction>SmartPolicy</misfire-instruction>
    <repeat-count>5</repeat-count>
    <repeat-interval>10000</repeat-interval>
  </simple>
</trigger>

```

```
</simple>  
</trigger>
```

Notice that the `<job-name>` and `<job-group>` elements refer to the same name and group that we specified in the job xml snippet above. These two elements are the means by which Quartz.Net ties the trigger to the job.

Executing our Custom Job

Now it is time to try executing our new custom job. We've already updated the `quartz_jobs.xml` file, but there is one step that we need to follow before we can have Quartz.Net schedule and execute our job. We need to move our `quickstart.dll` file to a directory where Quartz.Net can find it. Since Quartz.net is a .Net application, all of the rules that the CLR follows when trying to locate an assembly apply here. To make this as simple as possible, just copy the `quickstart.dll` file from your build folder to the Quartz.Net folder. If you're following the conventions we've been using so far, you'll want to copy the `quickstart.dll` file to `c:\QuartzNetBook\server`.

It's time to execute our new job. In the Quartz.Net command line prompt window you've got open, go ahead and start the server by typing `Quartz.Server.exe`. The Quartz.Net server will start up and after about a minute or so, both jobs will finish executing. You should now have two files in the `c:\QuartzNetBook\server` folder. `Run_times.txt` is the file that is created by the native job. It will have 5 new entries in it with the date and time, all about ten seconds apart. `QuickStartJob_runs.txt` is the file that was created by our new custom job and it should also have 5 entries in it, all about 10 seconds apart.

If this is not the case, look through the output of the scheduler to see what might have happened. If Quartz.Net was not able to locate your assembly or if there was an issue with the `quartz_jobs.xml` file, you will see one or several error messages in the console output.

Installing Quartz.Net as a Windows Service

So far, all of the examples that we have presented require you to start the Quartz.Net scheduler manually in a command prompt. For most production uses, this will not be acceptable and more than likely you will be installing Quartz.Net as a windows service. Fortunately, Quartz.Net supports this scenario without requiring you to do much.

To install Quartz.Net as a windows service, open a command prompt **as administrator** and change to the Quartz.Net directory: `c:\QuartzNetBook\server`. You should be in the same directory that has the `Quartz.Server.exe` file. The Quartz.Net server can be installed by typing the following in the command line:

```
Quartz.Server.exe install
```

If all goes well, you should see output similar to this in your window:



Figure 10

If your install completed successfully, then it's time to finish configuring the service in the service management console.

To open the service management console, type this in the command line:

```
services.msc
```

The services management window should open and it will be similar to the following:

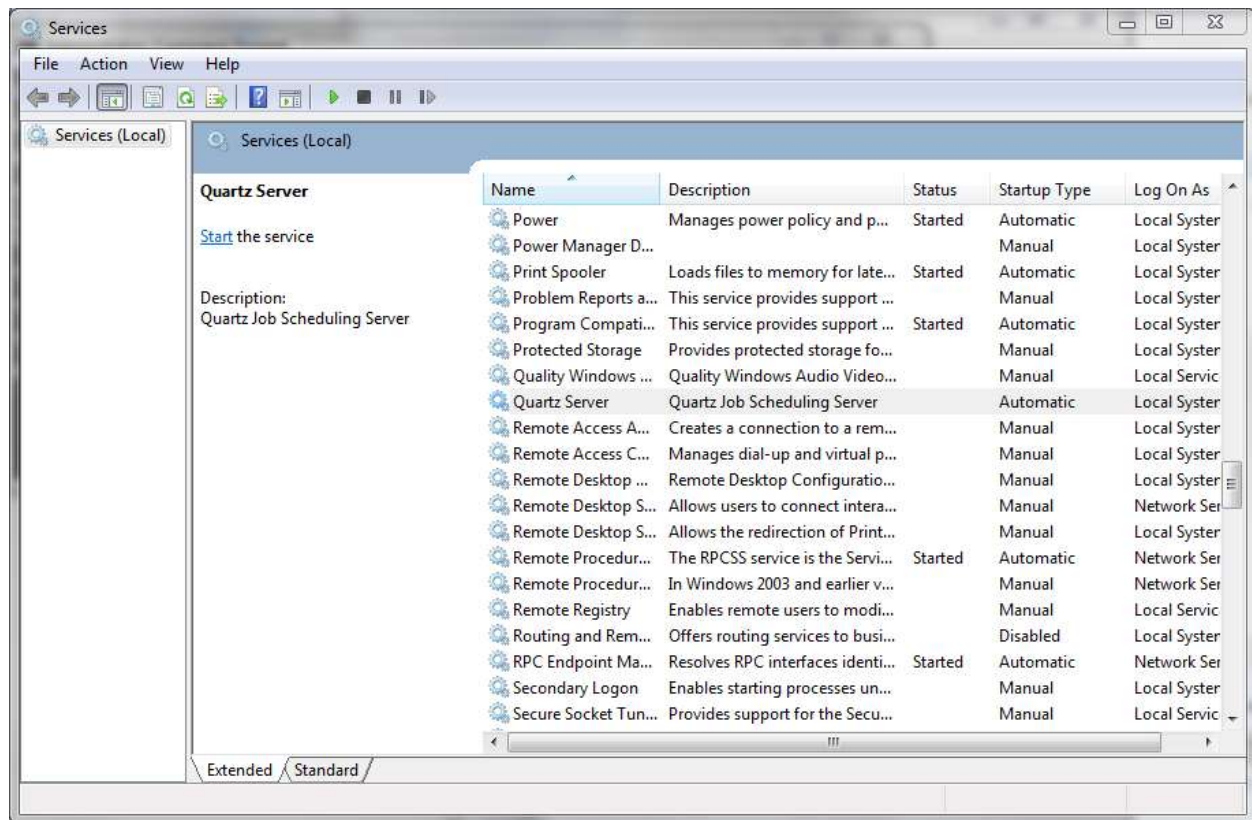


Figure 11

Locate the Quartz Server service and double click on it to edit it.

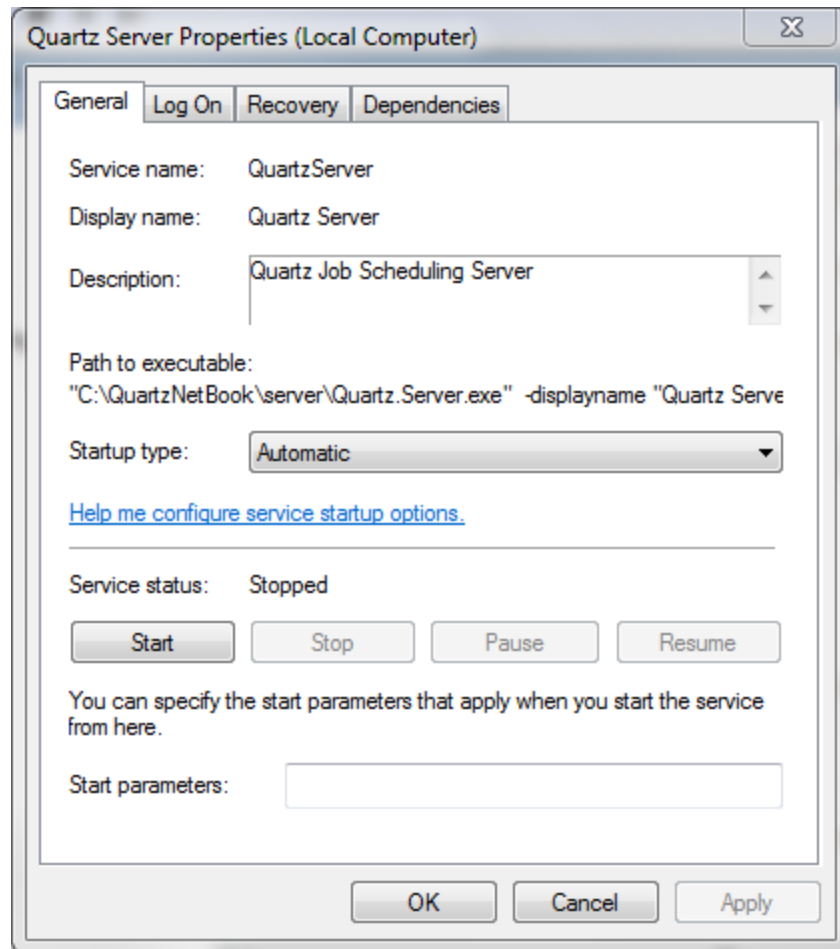


Figure 12

What you edit next depends on your requirements, but if you're just following the tutorial, you might want to change the startup type from automatic to manual, so that the server isn't running all the time. Additionally you might want to change the identity of the service if special permissions are required.

Uninstalling the Quartz.Net Service

To uninstall Quartz.Net as a windows service, open a command prompt as administrator and change to the Quartz.Net directory: `c:\QuartzNetBook\server`. You should be in the same directory that has the `Quartz.Server.exe` file. The Quartz.Net server can be uninstalled by typing the following in the command line:

```
Quartz.Server.exe uninstall
```

If all goes well, you should see output similar to this in your window:

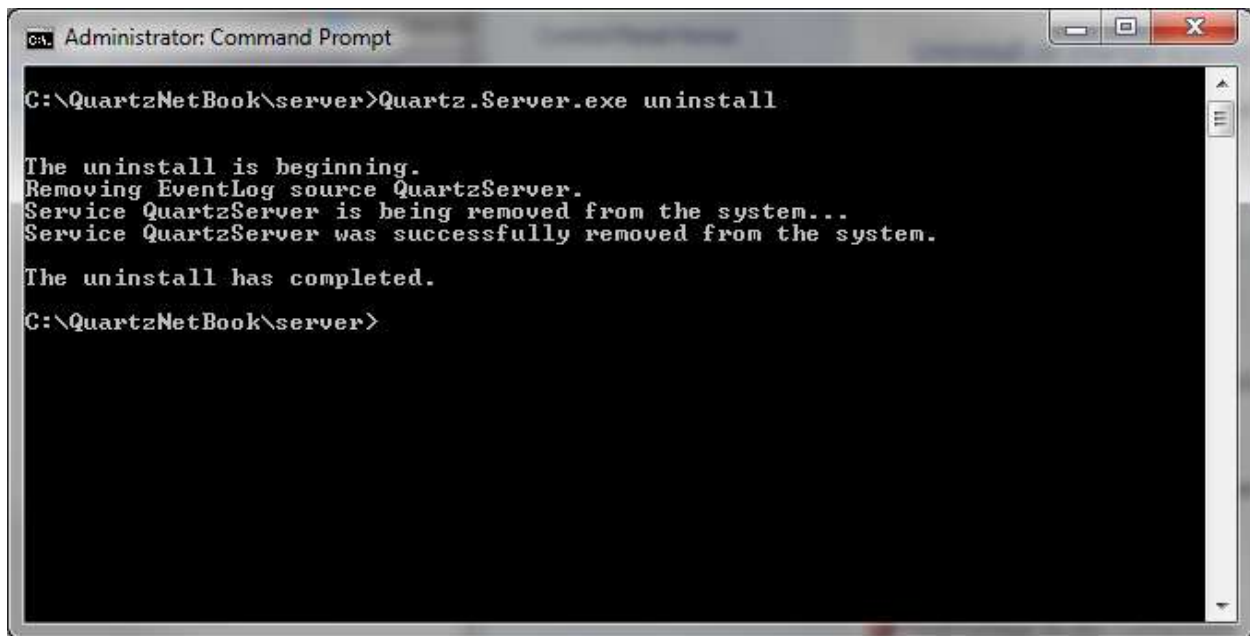


Figure 13

Embedding Quartz.Net in Your Application

So far we have seen that you can run Quartz.net directly from the command line. We also described how to install Quartz.Net as a windows service so that it is running all the time. Finally, we will describe how to embed Quartz.Net so that it runs within your application. While it's possible to embed Quartz.Net in an ASP.Net application, it's not recommended because web applications can be unceremoniously terminated by IIS. If you've embedded Quartz.Net in a web application and it gets abruptly terminated, well, bad things could happen.

So that we can illustrate the process of embedding Quartz.Net without spending too much time on the application itself, I've chosen to create a simple command line application that does nothing more than start the scheduler and wait until you press a key to terminate.

The code for this console application can be found under `c:\QuartzNetBook\samples\QuartzNetBookSamples`. Double click on the `QuartzNetBookSamples.sln` solution file to open it.

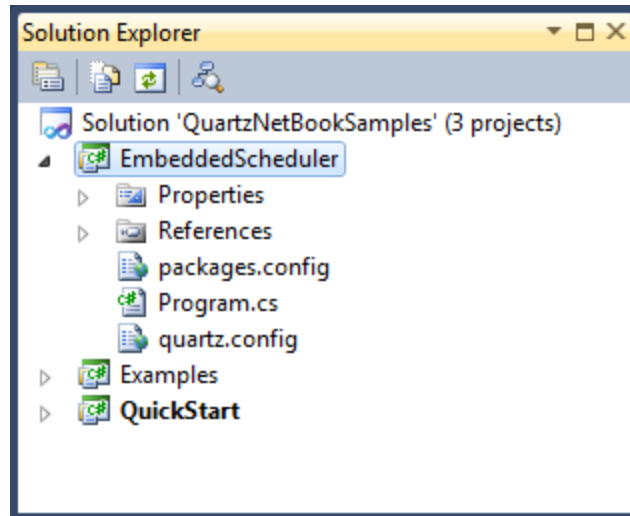


Figure 14

The EmbeddedScheduler project contains the code for our embedded scheduler. Let's open Program.cs and take a look inside. Here's the listing for the Main method of Program.cs.

```
1  static void Main(string[] args)
2  {
3      var factory = new StdSchedulerFactory();
4      var scheduler = factory.GetScheduler();
5      scheduler.Start();
6      while (!scheduler.IsStarted)
7      {
8          Console.WriteLine("Waiting for scheduler to start.");
9          Thread.Sleep(1000);
10     }
11     Console.WriteLine("IsStarted={0}", scheduler.IsStarted);
12     Console.WriteLine("InstanceId={0}", scheduler.SchedulerInstanceId);
13     Console.WriteLine("SchedulerName={0}", scheduler.SchedulerName);
14     Console.WriteLine("The scheduler is running. Press any key to stop");
15     Console.ReadKey();
16     Console.WriteLine("Shutting down scheduler");
17     scheduler.Shutdown(false);
18     while (!scheduler.IsShutdown)
19     {
20         Console.WriteLine("Waiting for scheduler to shutdown.");
21         Thread.Sleep(1000);
22     }
23     Console.WriteLine("IsShutdown={0}", scheduler.IsShutdown);
24     Console.WriteLine("The scheduler has been shutdown.");
25 }
```

Analyzing the Embedded Scheduler Program

Now we will look at the Program.cs listing above in detail.

Initializing and Starting the Scheduler

On line 3 we instantiate the standard scheduler factory. This factory is included in the Quartz.Net distribution and among other things, it is responsible for reading the configuration information for the scheduler. You may have noticed that there is a quartz.config file included in the EmbeddedScheduler project. Behind the scenes, the StdSchedulerFactory is reading this file and configuring the scheduler.

Next we ask the factory to create a scheduler for us and we store a reference to it. Storing a reference to the scheduler throughout the lifetime of your application is critical to make sure that the scheduler does not get garbage collected. If you do not hold a reference to the scheduler, it will stop running as soon as it goes out of scope.

After all this setup, we can now start the scheduler. This is shown on line 5. Lines 6-10 are just there to give the scheduler some time to start up, since it might not start up immediately. Lines 11-14 print out some scheduler information and then we stop and wait until the user presses a key. At this point the scheduler is fully operational and can be used by the application. In your application you will likely move on to your application's main screen or loop.

Shutting Down the Scheduler

When you're ready to stop the scheduler, you call the shutdown method on it, as shown on line 17. If you pass in true instead of false, the shutdown method will wait for all running jobs to complete. In the example we are asking the scheduler to terminate immediately.

Feedback

You can send any feedback, questions or comments on this draft version to jay@jayvilalta.com or follow the progress of this book at <http://jayvilalta.com/blog>